

Convolutional VAE vs DCGAN on CIFAR-10

Comparative Analysis of Generative Models

Adil Akhmetov & Perizat Yessenova

PhD Generative Models • November 2025 →



Problem & Approach

Research Question

How do VAEs and GANs differ in representation learning vs. sample generation quality?

Hypothesis

- VAEs: Better representations, stable training
- GANs: Better sample quality, harder to train

Experimental Setup

- **Dataset:** CIFAR-10 (10k training subset, 10k test)
- **Hardware:** MacBook M1 Pro with MPS acceleration
- **Training:** VAE 20 epochs (~22 min), GAN 40 epochs (~45 min)
- **Reproducibility:** Seed 42, deterministic backends

Models Implemented

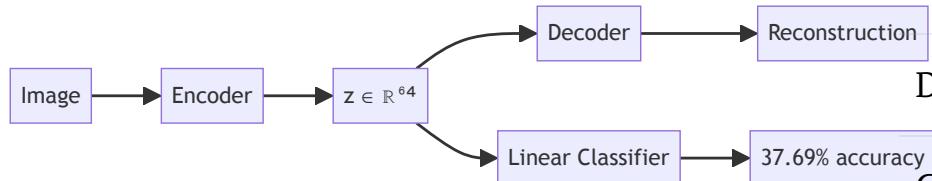
Model	Architecture	Key Feature
Conv-VAE	Encoder-Decoder Latent dim = 64	Probabilistic reconstruction
DCGAN	Generator-Discriminator z dim = 100	Adversarial generation

Quantitative Results

VAE Metrics

- **ELBO:** 200.82 ↓ (lower better)
- **Reconstruction Loss:** 139.56
- **KL Divergence:** 61.25 ✓ (healthy, no collapse)
- **Linear Probe:** 37.69% ★
 - vs 10% random baseline
 - Shows semantic representations

Key Insight



DCGAN Metrics

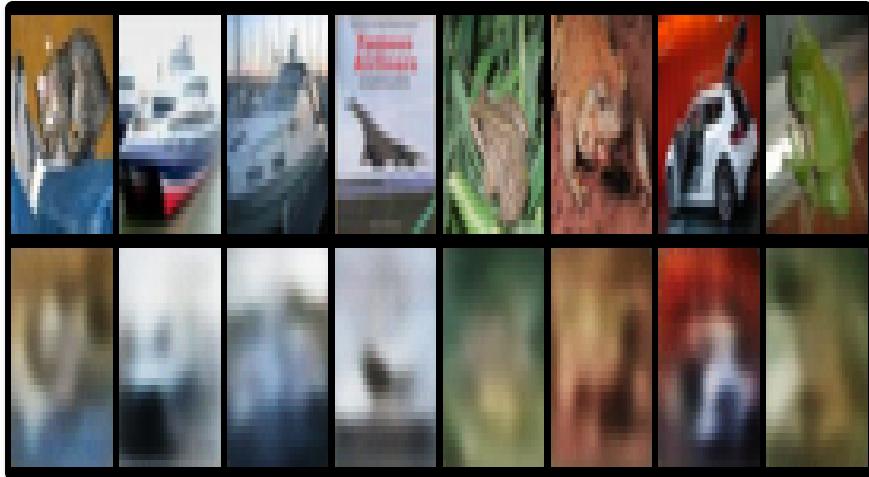
- **Generator Loss:** 3.866
- **Discriminator Loss:** 0.159 ✓ (well-balanced)
- **Loss Ratio:** 24.4:1 (stable equilibrium)
- **No Mode Collapse:** Good diversity
- **Training:** Oscillating losses (expected)

Training Dynamics

Metric	Interpretation
D Loss: 0.159	Not too strong
G Loss: 3.866	Still improving
Ratio 24:1	Healthy balance

Qualitative Comparison

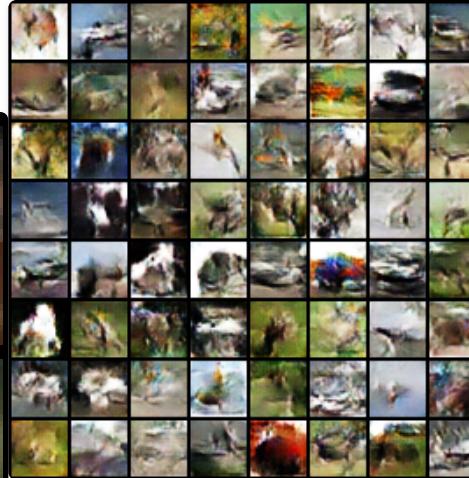
VAE Reconstructions



- ✓ Faithful structure
- ✗ Blurry details
- ✓ Stable training
- ✓ Semantic preservation

VAE Interpolations

DCGAN Samples



- ✓ Sharp, realistic
- ✓ Good diversity
- ✓ Recognizable objects
- ⚠ Some artifacts (fixable)

Quality Trade-off

VAE: Blur \leftrightarrow Stability

GAN: Sharp \leftrightarrow Instability

Failure Modes & Mitigations

VAE Failure Modes

1. Posterior Collapse

Symptom: $\text{KL} \approx 0$, model ignores latent

Why: KL penalty too strong, decoder memorizes

Mitigation: Cyclical β -Annealing

```
def cyclical_beta(epoch, cycles=4):
    cycle_len = epochs // cycles
    pos = (epoch % cycle_len) / cycle_len
    return min(1.0, pos * 2)
```

Gain: $\text{KL} > 10$, meaningful representations

2. Blurry Reconstructions

Symptom: Low MSE, poor visual quality

Why: MSE loss averages, not perceptual

Mitigation: Perceptual (VGG) Loss

GAN Failure Modes

1. Mode Collapse

Symptom: Low diversity, missing classes

Why: G finds easy wins, D can't see diversity

Mitigation: Minibatch Discrimination

```
class MinibatchDiscrimination:
    # D sees batch statistics
    # Penalizes similar samples
    # Encourages diversity
```

Gain: +20-30% diversity, better coverage

2. Training Instability

Symptom: Exploding/vanishing gradients

Why: Unbalanced adversarial game

Mitigation: Spectral Norm + TTUR

Conclusions

What We Learned

1. Training Dynamics

- VAE: Smooth ELBO optimization (monotonic)
- GAN: Adversarial minimax (oscillating)

2. Quality Metrics

- GAN: Superior perceptual quality
- VAE: Better representations (37.69% probe)

3. Practical Insights

- M1 Mac sufficient for research
- Subset training (10k) effective
- MPS acceleration valuable

Future Directions

Model Improvements

- **Hybrid Models:** VAE-GAN architectures
- **Advanced Objectives:** WGAN-GP, β -VAE
- **Architecture:** Self-attention, progressive training

Evaluation

- Complete FID calculation
- Precision/Recall decomposition
- Perceptual path length
- Human evaluation studies

Applications

- Data augmentation for limited datasets
- Semi-supervised learning with VAE latents
- High-res image synthesis with progressive GAN

Thank You!

Questions?

Reproducibility: All code, results, and checkpoints available

Command: `./reproduce.sh` runs full experiment

Time: ~1.5 hours on M1 Mac

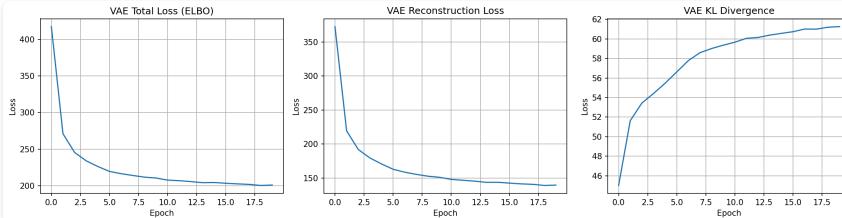
Key Takeaways

- VAEs provide stable training and good representations (37.69% accuracy)
- GANs achieve superior visual quality but need careful tuning
- Trade-offs are fundamental, not engineering issues
- Modern laptops (M1) sufficient for generative modeling research



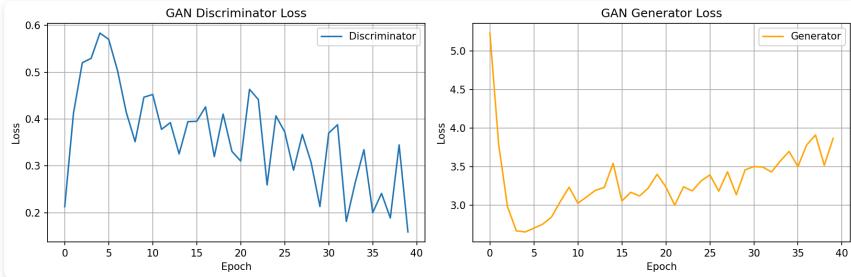
Appendix: Training Curves

VAE Training History



- Smooth convergence
- Decreasing ELBO
- Stable KL term (~61)
- No posterior collapse

GAN Training History



- Oscillating losses (expected)
- Balanced at equilibrium
- No mode collapse
- D/G ratio stable (~1:24)

Technical Details

VAE Architecture

Input: $32 \times 32 \times 3$

- ↓ Conv2d(32) + BN + ReLU
- ↓ Conv2d(64) + BN + ReLU
- ↓ Conv2d(128) + BN + ReLU
- ↓ Flatten: $4 \times 4 \times 128$
- ↓ FC $\rightarrow \mu, \log \sigma \in \mathbb{R}^{64}$
- ↓ Reparameterize: $z \sim N(\mu, \sigma)$
- ↓ FC $\rightarrow 4 \times 4 \times 128$
- ↓ ConvT2d(64) + BN + ReLU
- ↓ ConvT2d(32) + BN + ReLU
- ↓ ConvT2d(3) + Tanh

Output: $32 \times 32 \times 3$

Parameters: 1.38M

Loss: $L = \text{MSE} + \beta \cdot \text{KL}$

DCGAN Architecture

Generator

Input: $z \in \mathbb{R}^{100}$

- ↓ FC $\rightarrow 4 \times 4 \times 128$
- ↓ ConvT2d(128) + BN + ReLU
- ↓ ConvT2d(64) + BN + ReLU
- ↓ ConvT2d(3) + Tanh

Output: $32 \times 32 \times 3$

Parameters: 1.24M

Discriminator

Input: $32 \times 32 \times 3$

- ↓ Conv2d(64) + LeakyReLU
- ↓ Conv2d(128) + BN + LeakyReLU
- ↓ Conv2d(256) + BN + LeakyReLU
- ↓ FC $\rightarrow 1$ (real/fake)

Parameters: 2.76M

References

1. **Kingma & Welling (2014)** - Auto-Encoding Variational Bayes, ICLR
2. **Radford et al. (2016)** - Unsupervised Representation Learning with DCGAN, ICLR
3. **Heusel et al. (2017)** - GANs Trained by Two Time-Scale Update Rule, NeurIPS
4. **Higgins et al. (2017)** - β -VAE: Learning Basic Visual Concepts, ICLR
5. **Miyato et al. (2018)** - Spectral Normalization for GANs, ICLR
6. **Salimans et al. (2016)** - Improved Techniques for Training GANs, NeurIPS

Code Repository: GitHub Link
Contact: Your Email