

Project 2 (By two person)

You should implement unfolded Algorithms, Provide Visuals, Provide as possible graphical results

Below is a step-by-step guide on how to train a reinforcement learning agent (e.g., using **SAC**, **DDPG**, **PPO**, or **TD3**) with [**Gym-PyBullet-Drones**](#). This guide assumes you're using **Python 3.x** and will leverage **Stable Baselines3** for the RL algorithms.

1. Install the Required Packages

1. Create/Activate a Python Virtual Environment (*recommended*):

```
python -m venv venv
source venv/bin/activate  # On Linux/Mac
# or
venv\Scripts\activate.bat  # On Windows
```

2. Install the necessary libraries:

```
pip install --upgrade pip
pip install pybullet
pip install stable-baselines3
pip install git+https://github.com/utiasDSL/gym-pybullet-drones.git
```

This will install:

- **pybullet** (physics simulation)
- **stable-baselines3** (popular RL algorithms like SAC, TD3, PPO, DDPG)
- **gym-pybullet-drones** (the custom Gym environments for drones)

2. Choose an Environment

Gym-PyBullet-Drones provides several environments (or “aviaries”), for example:

- **hover-aviary-v0**: Single-drone environment where the goal is to hover at a specified position.

- takeoff-aviary-v0: Single-drone environment where the drone must take off and reach a target height.
- flythrugate-aviary-v0: Single or multi-drone environment for flying through gates (a more complex task).
- tune-aviary-v0: An environment specifically designed for classical (PID) control tuning.

For simplicity, we'll use hover-aviary-v0 in the example below.

3. Minimal Example: Training With SAC

Here is a short Python script (`train_sac.py`) illustrating how to train **SAC** on hover-aviary-v0:

```
import gym
import gym_pybullet_drones # This registers the drone environments
from stable_baselines3 import SAC
from stable_baselines3.sac.policies import MlpPolicy

def main():
    # 1. Create the environment
    env = gym.make("hover-aviary-v0")

    # 2. Instantiate the RL model (SAC in this case)
    model = SAC(
        policy=MlpPolicy,
        env=env,
        verbose=1,          # Set to 1 (info) or 2 (debug)
        learning_rate=1e-3,
        buffer_size=100000,
        batch_size=256,
        gamma=0.99,
        tau=0.02
    )

    # 3. Train the agent
    timesteps = 50000
    model.learn(total_timesteps=timesteps)

    # 4. Save the model
    model.save("sac_hover_aviary")

    # Optional: close the environment if needed
    env.close()

if __name__ == "__main__":
    main()
```

Key Points

- **Policy:** MlpPolicy means a Multi-Layer Perceptron (fully-connected neural net).
- **Hyperparameters:** The above defaults often work decently, but you may want to tune learning_rate, batch_size, etc., based on performance.
- **Timesteps:** Adjust 50000 (or more) depending on your computational resources and how complex the task is.

4. Training With Other Algorithms

In **Stable Baselines3**, you can replace SAC with other continuous-control algorithms such as **DDPG**, **TD3**, or even a discrete-friendly algorithm like **PPO** (though PPO is widely used in continuous action spaces as well). The code structure is the same; only the import changes:

- **DDPG:**
- **TD3:**
- **PPO** (though typically used for discrete or continuous):

Everything else (environment creation, saving, etc.) remains basically the same.

5. Monitoring Training & Plotting

1. Logging:

- By default, Stable Baselines3 logs to the console.
- You can optionally pass a tensorboard_log parameter to log data for TensorBoard:

```
model = SAC("MlpPolicy", env, verbose=1, tensorboard_log="./tensorboard/")
model.learn(total_timesteps=50000)
```

- Then, run tensorboard --logdir ./tensorboard/ to visualize training curves in your browser.

2. Plotting:

- You can record rewards by wrapping your environment or by using callback functions.
- See [Stable Baselines3's documentation](#) on how to implement callbacks and track custom metrics.

6. Evaluating the Trained Model

After training, you likely want to see how the drone performs. Here's a quick script (`evaluate_sac.py`) to load the model and visualize the drone in PyBullet:

```
import gym
import gym_pybullet_drones
from stable_baselines3 import SAC

def evaluate_model():
    # 1. Load the environment
    env = gym.make("hover-aviary-v0", gui=True)  # gui=True to see PyBullet
    UI

    # 2. Load the trained model
    model = SAC.load("sac_hover_aviary", env=env)

    # 3. Run episodes
    obs = env.reset()
    for _ in range(1000):  # 1000 steps
        # Get the action from the model
        action, _states = model.predict(obs, deterministic=True)
        obs, reward, done, info = env.step(action)

        # Optionally, you can print or log the reward
        # print("Reward:", reward)

        if done:
            obs = env.reset()

    env.close()

if __name__ == "__main__":
    evaluate_model()
```

Key Points

- `gui=True`: In `gym-pybullet-drones`, specifying `gui=True` will render the environment so you can visually watch the drone.
- `model.predict(obs, deterministic=True)`: By default, SAC is stochastic, so setting `deterministic=True` means you see the “greedy” (exploiting) version of the policy.

7. Tips for Successful Training

1. Simplify the Task:

If the drone is too unstable:

- Start with an easier environment like takeoff-aviary-v0 or a smaller action range.
- Use a well-shaped reward function (e.g., negative distance to a hover point, plus penalties for excessive tilt).

2. Hyperparameter Tuning:

- Learning rate, batch_size, buffer_size, and exploration strategies (for DDPG/TD3) can dramatically affect training.
- Start with default Stable Baselines3 hyperparameters, then adjust systematically.

3. Check Observations & Actions:

- Make sure you understand the observation space (state vector) and the action space (motor thrusts, velocity commands, or orientation) provided by the environment.
- If your agent never learns, you might need to modify these or re-scale the action space.

4. Logging & Debugging:

- Inspect reward logs frequently. If rewards don't improve or remain negative, adjust your environment or reward shaping.

5. Speeding Up Training:

- Increase the n_envs for parallel environments if your hardware supports it.
- Use GPU acceleration for the policy network if possible.

8. Next Steps & Extensions

- **Multiple Waypoints:** Extend the environment to require the drone to pass through multiple waypoints.
- **Obstacle Avoidance:** Add static or moving objects in the environment with collision penalties.
- **Domain Randomization:** Randomize gravity, wind, or drone mass to encourage robust policies.
- **Multi-Agent:** Gym-PyBullet-Drones supports multi-drone setups—train multi-agent RL with cooperation or competition tasks.

Final Remarks

That's it! With these instructions, you should be able to:

1. Install and set up Gym-PyBullet-Drones.
2. Train with an RL algorithm of your choice (SAC, TD3, DDPG, PPO, etc.) using Stable Baselines3.
3. Evaluate and visualize the learned drone behavior in PyBullet.

This workflow is a great starting point for experimenting with more complex drone control tasks and reward structures. Good luck with your training, and feel free to ask if you have any more questions!

**Or run and train Bipedal waslking robot in
Open AI gym**

Then you will get maximum 80 points

