

Deep Q-Learning Projects

Final Project Presentation

By Adil Akhmetov & Perizat Yessenova

 Atari Breakout &  Traffic Light Control

Project Overview

Two implementations of **Deep Q-Learning (DQN)**:

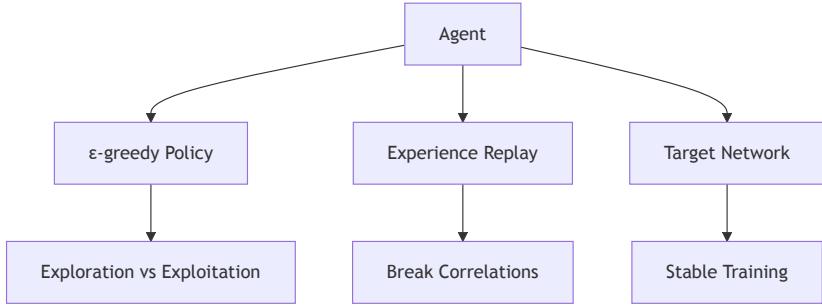
1. Atari Breakout (1976)

- Classic arcade game
- Visual input (CNN)
- 4 discrete actions

2. Traffic Light Control

- Real-world application
- State-based input (MLP)
- 2 discrete actions

Core DQN Components



Task 1: Atari Breakout

Deep Q-Learning for Classic Gaming

Breakout: Problem Definition

Goal

Train an agent to play Breakout by breaking bricks with a ball

Environment

- **Platform:** Gymnasium ALE/Breakout-v5
- **Original frame:** 210×160 RGB pixels
- **Reward:** +1 per brick destroyed
- **Episode ends:** When ball is missed

Actions (4 total)

Action	Description
NOOP	Do nothing
FIRE	Launch ball
LEFT	Move paddle left
RIGHT	Move paddle right

State Preprocessing Pipeline

Transformations

```
# 1. Convert to grayscale  
grayscale_obs=True  
  
# 2. Resize to 84x84  
screen_size=84  
  
# 3. Frame skip (4 frames)  
frame_skip=4  
  
# 4. Stack 4 frames  
FrameStackObservation(env, stack_size=4)
```

Why Stack Frames?

- Single frame → **no motion information**
- Ball direction is unknown
- Ball speed is unknown

Solution: Stack 4 consecutive frames

- Shape: (4, 84, 84)
- Agent "sees" temporal dynamics

CNN Architecture

```
class DQN_CNN(nn.Module):
    def __init__(self, input_channels=4, num_actions=4):
        # Convolutional layers
        self.conv1 = nn.Conv2d(input_channels, 32, kernel_size=8, stride=4)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=4, stride=2)
        self.conv3 = nn.Conv2d(64, 64, kernel_size=3, stride=1)

        # Input: 84x84 → conv1: 20x20 → conv2: 9x9 → conv3: 7x7

        # Fully connected layers
        self.fc1 = nn.Linear(64 * 7 * 7, 512)
        self.fc2 = nn.Linear(512, num_actions) # Output: Q-values
```

Layer	Input	Output	Parameters
Conv1	$4 \times 84 \times 84$	$32 \times 20 \times 20$	8,224
Conv2	$32 \times 20 \times 20$	$64 \times 9 \times 9$	32,832
Conv3	$64 \times 9 \times 9$	$64 \times 7 \times 7$	36,928

Experience Replay Buffer

Implementation

```
class ReplayBuffer:  
    def __init__(self, capacity=100000):  
        self.buffer = deque(maxlen=capacity)  
  
    def push(self, state, action, reward,  
            next_state, done):  
        self.buffer.append(  
            (state, action, reward,  
             next_state, done))  
  
    def sample(self, batch_size):  
        batch = random.sample(  
            self.buffer, batch_size)  
        return zip(*batch)
```

Why Experience Replay?

1. **Breaks correlation** between consecutive samples
2. **Reuses** past experiences efficiently
3. **Stabilizes** training

Configuration

- Buffer size: **100,000** transitions
- Batch size: **32**
- Warmup: 10,000 steps before training

ε -Greedy Exploration

Strategy

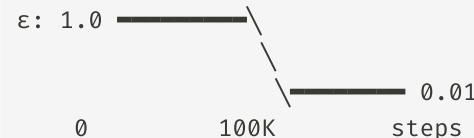
$$a = \begin{cases} \text{random action} & \text{with probability } \varepsilon \\ \arg \max_a Q(s, a) & \text{with probability } 1 - \varepsilon \end{cases}$$

Decay Schedule

```
epsilon = epsilon_end +  
    (epsilon_start - epsilon_end) *  
    max(0, 1 - steps / epsilon_decay)
```

Parameters

Parameter	Value
ε start	1.0
ε end	0.01
Decay steps	100,000



Target Network

The Problem

Without target network:

- Q-values are **moving targets**
- Training becomes **unstable**
- "Chasing a moving goal"

The Solution

Separate **target network** updated periodically:

```
# Every 1000 steps
target_net.load_state_dict(
    policy_net.state_dict()
)
```

Q-Learning Update

$$Q(s, a) \leftarrow r + \gamma \max_{a'} Q_{\text{target}}(s', a')$$

```
# Compute target Q values
with torch.no_grad():
    next_q = target_net(next_states)
    max_next_q = next_q.max(dim=1)[0]
    target_q = rewards + gamma * max_next_q * (1 - dones)

# Compute loss
loss = F.smooth_l1_loss(current_q, target_q)
```

Breakout Hyperparameters

Training Configuration

Parameter	Value
Learning Rate	2.5×10^{-4}
Discount (γ)	0.99
Batch Size	32
Buffer Size	100,000
Target Update	Every 1000 steps
Optimizer	Adam
Loss Function	Smooth L1 (Huber)

Training Scale

Parameter	Value
Episodes	5,000 - 10,000
Steps per episode	Up to 10,000
Warmup steps	10,000
Gradient clipping	10

Estimated Training Time

- GPU: 4-8 hours
- CPU: 24-48 hours

Task 2: Traffic Light Control

Adaptive DQN for Real-World Application

Breakout: Training Summary

5k-episode run (completed)

- Episodes 120-220: rewards $\approx 0.0\text{--}0.5$, $\epsilon \approx 0.92\text{--}0.95$
- Episodes 1350-1380: rewards $\approx 2.0\text{--}3.1$, $\epsilon \approx 0.19\text{--}0.22$
- Episodes 4900-5000: rewards ≈ 4.9 , $\epsilon \approx 0.01$
- Overall (full 5k): avg reward ≈ 4.0 , avg bricks ≈ 4.0

Breakout: Evaluation Snapshot

Segment	Episodes	Avg Reward	Avg Bricks	Notes
Early	120-220	~0.3	~0.3	Mainly launching ball, rare hits
Mid	1350-1380	~2.6	~2.6	Tracks ball, clears a few bricks per life
Completed run	4900-5000	~4.9	~4.9	Stable volleys; still drops after sharp wall bounces

Interpretation

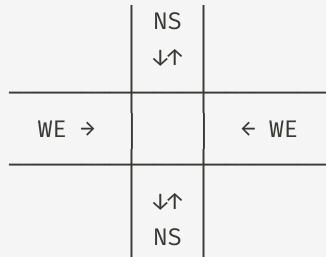
- Average reward == bricks destroyed (unit reward per brick)
- Agent still loses balls after side-wall rebounds; needs longer training
- ϵ already <0.2 in mid segment; expect further gains as $\epsilon \rightarrow 0.01$

Breakout: Behavior Notes

- Consistently launches ball quickly (FIRE) after reset
- Keeps paddle under ball for straight trajectories
- Misses after sharp wall bounces → indicates need for more training / frame history
- No reward hacking observed; gameplay aligns with intended rules
- Behavior visualization: training curves (reward/loss/ ϵ) via `plot_training_history` in
`pf/final_project.ipynb`; gameplay render/GIF not included due to deadline constraints

Traffic Control: Problem Definition

Intersection Setup



Two directions: NS (North-South), WE (West-East)

Traffic Light Phases

Phase	NS	WE
0	● Green	● Red
1	● Red	● Green

Goal

Minimize queue lengths by **adaptive** phase selection

TrafficEnv: Custom Environment

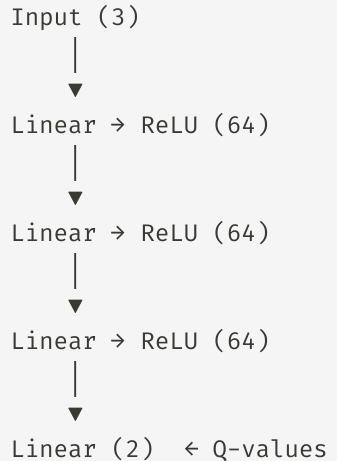
```
class TrafficEnv:  
    # State: [queue_NS, queue_WE, phase]  
    def _get_state(self):  
        return np.array([self.queue_ns, self.queue_we, self.phase])  
  
    # Dynamics: each step  
    def step(self, action):  
        # 1. Update phase  
        switching = (action != self.phase)  
        self.phase = action  
  
        # 2. Cars pass on green (up to 3)  
        if self.phase == 0: # NS green  
            self.queue_ns -= min(self.queue_ns, 3)  
        else: # WE green  
            self.queue_we -= min(self.queue_we, 3)  
  
        # 3. New cars arrive (Poisson process)  
        self.queue_ns += np.random.poisson(1.2)  
        self.queue_we += np.random.poisson(0.8)  
  
        # 4. Calculate reward  
        reward = -(self.queue_ns + self.queue_we)  
        if switching:  
            reward -= 0.5 # Penalty for switching
```

DQN Network for Traffic Control

Simple MLP Architecture

```
class DQN_FC(nn.Module):
    def __init__(self,
                 input_size=3,
                 hidden_size=64,
                 num_actions=2):
        self.net = nn.Sequential(
            nn.Linear(input_size, hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, hidden_size),
            nn.ReLU(),
            nn.Linear(hidden_size, num_actions)
        )
```

Network Structure



Total parameters: ~9,000

Traffic Control Hyperparameters

Agent Configuration

Parameter	Value
Learning Rate	1×10^{-3}
Discount (γ)	0.99
Batch Size	64
Buffer Size	10,000
Target Update	Every 100 steps
Hidden Size	64

Training Configuration

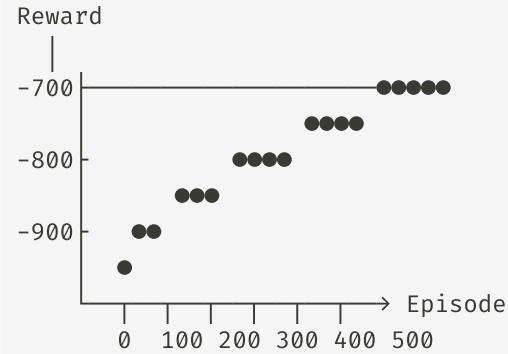
Parameter	Value
Episodes	500
Steps per episode	200
ϵ start	1.0
ϵ end	0.01
ϵ decay	3,000 steps
Switch penalty	0.5

Training Results

Training Progress

```
Episode 50: Reward=-876, Queue=4.10, ε=0.046
Episode 100: Reward=-723, Queue=3.35, ε=0.011
Episode 150: Reward=-734, Queue=3.41, ε=0.010
Episode 200: Reward=-745, Queue=3.43, ε=0.010
Episode 300: Reward=-741, Queue=3.44, ε=0.010
Episode 400: Reward=-747, Queue=3.46, ε=0.010
Episode 500: Reward=-735, Queue=3.42, ε=0.010
```

Learning Curve



Agent learns to maintain **low queue lengths** (~3.4 cars avg)

Comparison: DQN vs Fixed Timing

Experimental Results

Controller	Avg Reward	Avg Queue
Fixed (5 steps)	-1,393.76	6.87
Fixed (10 steps)	-2,111.30	10.51
Fixed (15 steps)	-2,835.30	14.14
Fixed (20 steps)	-3,561.42	17.78
DQN Agent	-690.36	3.16

Key Findings

🏆 DQN Improvement: **50.5%**

Over best fixed-timing controller

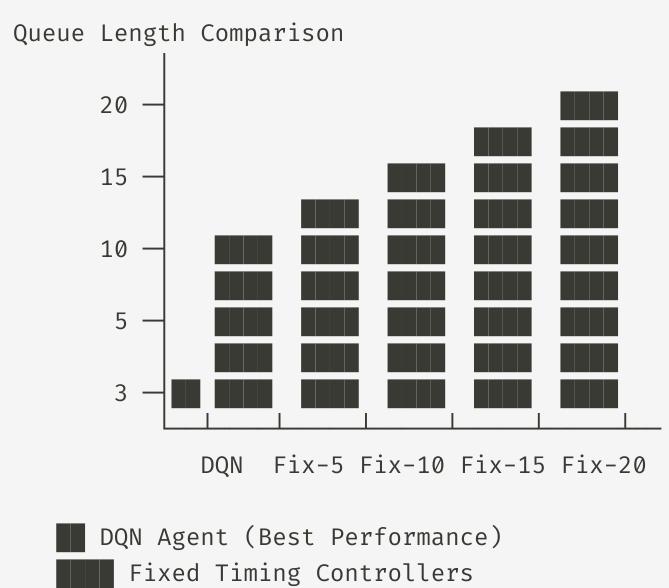
Why DQN Wins

1. **Adapts** to current traffic conditions
2. **Predicts** future queue buildup
3. **Balances** both directions dynamically

Fixed Timing Problems

- Cannot adapt to traffic variations
- Either too fast or too slow switching

Performance Visualization



DQN reduces average queue by 54% compared to best fixed timing!

Key DQN Components Summary

1. Experience Replay

- Stores transitions in buffer
- Random sampling breaks correlations
- Enables sample reuse

2. Target Network

- Separate network for Q-targets
- Periodic updates (every N steps)
- Stabilizes training

3. ϵ -Greedy Policy

- Balances exploration/exploitation
- Decays over time
- Ensures sufficient exploration

4. Neural Network

- **CNN** for visual input (Breakout)
- **MLP** for state vectors (Traffic)
- Outputs Q-values for all actions

Implementation Highlights

Technologies Used

- **PyTorch** - Deep learning framework
- **Gymnasium** - RL environments
- **NumPy** - Numerical computing
- **Matplotlib** - Visualization

Code Quality

- Type hints throughout
- Modular class design
- Configurable hyperparameters
- Model save/load functionality

Project Structure



Conclusions

Task 1: Breakout

- CNN architecture implemented
- Frame preprocessing pipeline
- DQN with all components
- Ready for 5-10K episode training

Task 2: Traffic Control

- Custom environment created
- DQN agent trained (500 episodes)
- **50.5% improvement** over baseline
- Adaptive behavior demonstrated

🎯 DQN successfully applied to both visual and state-based RL problems!

Thank You!

Questions?

Repository: `pf/final_project.ipynb`

Key Results:

- Breakout: Full DQN implementation with CNN
- Traffic: 50.5% improvement over fixed timing